

Logistic Regression

3 Jun 2019

Vineeth N Balasubramanian



आई आई टी हैदराबाद
IIT Hyderabad

Classification Methods

- k-Nearest Neighbors
- Decision Trees
- Naïve Bayes
- Support Vector Machines
- **Logistic Regression**
- Neural Networks
- Ensemble Methods (Boosting, Random Forests)

Regression Formulation

Given x , want to predict an estimate \hat{y} of y , which minimizes the discrepancy (L) between \hat{y} and y .


$$\begin{aligned} L(\hat{y}; y) &:= |\hat{y} - y| && \text{Absolute error} \\ \text{Loss} &:= (\hat{y} - y)^2 && \text{Squared error} \end{aligned}$$

A **linear predictor** f , can be defined by the slope w and the intercept w_0 :

$$\hat{f}(\vec{x}) := \vec{w} \cdot \vec{x} + w_0$$

which minimizes the prediction loss.

$$\min_{w, w_0} \mathbb{E}_{\vec{x}, y} [L(\hat{f}(\vec{x}), y)]$$




Slide Credit: Nakul Verma, Columbia Univ

Parametric (vs) Non- parametric Regression

If we assume a particular form of the regressor:

Parametric regression


Goal: to learn the parameters which yield the minimum error/loss



If no specific form of regressor is assumed:

Non-parametric regression

Goal: to learn the predictor directly from the input data that yields the minimum error/loss



Slide Credit: Nakul Verma, Columbia Univ

Linear Regression

Want to find a **linear predictor** f , i.e., w (intercept w_0 absorbed via lifting):

$$\hat{f}(\vec{x}) := \vec{w} \cdot \vec{x}$$


which minimizes the prediction loss over the population.

$$\min_{\vec{w}} \mathbb{E}_{\vec{x}, y} [L(\hat{f}(\vec{x}), y)]$$

We estimate the parameters by minimizing the corresponding loss on the training data:

$$\begin{aligned} \arg \min_w \frac{1}{n} \sum_{i=1}^n [L(\vec{w} \cdot \vec{x}_i, y_i)] \\ = \arg \min_w \frac{1}{n} \sum_{i=1}^n (\vec{w} \cdot \vec{x}_i - y_i)^2 \end{aligned}$$

for squared error



Slide Credit: Nakul Verma, Columbia Univ

Linear Regression

Linear predictor with squared loss:

$$\arg \min_w \frac{1}{n} \sum_{i=1}^n (\vec{w} \cdot \vec{x}_i - y_i)^2$$

$$= \arg \min_w \left\| \begin{pmatrix} \dots \mathbf{x}_1 \dots \\ \dots \mathbf{x}_i \dots \\ \dots \mathbf{x}_n \dots \end{pmatrix} \begin{pmatrix} \mathbf{w} \end{pmatrix} - \begin{pmatrix} y_1 \\ y_i \\ y_n \end{pmatrix} \right\|^2$$

$$= \arg \min_w \|X\vec{w} - \vec{y}\|_2^2$$

Unconstrained problem!

Can take the gradient and examine the stationary points!

Why need not check the second order conditions?

Slide Credit: Nakul Verma, Columbia Univ

Linear Regression

Best fitting w :

$$\frac{\partial}{\partial \vec{w}} \|X\vec{w} - \vec{y}\|^2 = 2X^\top(X\vec{w} - \vec{y})$$

$$X^\top X\vec{w} = X^\top \vec{y} \quad \text{At a stationary point}$$

$$\Rightarrow \vec{w}_{\text{ols}} = (X^\top X)^\dagger X^\top \vec{y}$$

Pseudo-inverse

Also called the Ordinary Least Squares (OLS)

The solution is unique and stable when $X^\top X$ is invertible

Slide Credit: Nakul Verma, Columbia Univ

Regularized Least-Squared Regression

- Complex models (lots of parameters) often prone to overfitting.
- Overfitting can be reduced by imposing a constraint on the overall magnitude of the parameters.
- Two common types of regularization in linear regression:

- **L₂ regularization (a.k.a. ridge regression):** Find \mathbf{w} which minimizes:

$$\sum_{j=1}^N (y_j - \sum_{i=0}^d w_i \cdot x_i)^2 + \lambda \sum_{i=1}^d w_i^2$$

- λ is the regularization parameter: bigger λ imposes more constraint

- **L₁ regularization (a.k.a. lasso):** Find \mathbf{w} which minimizes:

$$\sum_{j=1}^N (y_j - \sum_{i=0}^d w_i \cdot x_i)^2 + \lambda \sum_{i=1}^d |w_i|$$


k-NN Regression

- Calculate the mean value of the k nearest training examples rather than calculate their most common value

$$f : \mathbb{R}^d \rightarrow \mathbb{R} \qquad \hat{f}(x_q) \leftarrow \frac{\sum_{i=1}^k f(x_i)}{k}$$

Logistic Regression

- To predict an outcome variable that is categorical from one or more categorical or continuous predictor variables.
- Used because having a categorical outcome variable violates the assumption of linearity in normal regression.
- Let X be the data instance, and Y be the class label: Learn $P(Y|X)$ directly
 - Let $W = (W_1, W_2, \dots, W_n)$, $X = (X_1, X_2, \dots, X_n)$, $\mathbf{W} \cdot \mathbf{X}$ is the dot product
 - Sigmoid function:
$$P(Y = 1 | \mathbf{X}) = \frac{1}{1 + e^{-\mathbf{W} \cdot \mathbf{X}}}$$



Logistic Regression

- Generative or Discriminative?

Logistic Regression

- Generative classifier, e.g., Naïve Bayes:
 - Assume some functional form for **$P(\mathbf{X}|\mathbf{Y})$** , **$P(\mathbf{Y})$**
 - Estimate parameters of $P(\mathbf{X}|\mathbf{Y})$, $P(\mathbf{Y})$ directly from training data
 - Use Bayes rule to calculate $P(\mathbf{Y}|\mathbf{X}=\mathbf{x})$
 - This is 'generative' model
 - Indirect computation of $P(\mathbf{Y}|\mathbf{X})$ through Bayes rule
 - But, can generate a sample of the data, $P(\mathbf{X}) = \sum_y P(y)P(\mathbf{X} | y)$
- Discriminative classifier, e.g., Logistic Regression:
 - Assume some functional form for **$P(\mathbf{Y}|\mathbf{X})$**
 - Estimate parameters of $P(\mathbf{Y}|\mathbf{X})$ directly from training data
 - This is the 'discriminative' model
 - Directly learn $P(\mathbf{Y}|\mathbf{X})$

Logistic Regression

- In logistic regression, we learn the conditional distribution $P(y|x)$
- Let $p_y(x;w)$ be our estimate of $P(y|x)$, where w is a vector of adjustable parameters.
- Assume there are two classes, $y = 0$ and $y = 1$ and

$$p_1(\mathbf{x}; \mathbf{w}) = \frac{1}{1 + e^{-\mathbf{w}\mathbf{x}}} \qquad p_0(\mathbf{x}; \mathbf{w}) = 1 - \frac{1}{1 + e^{-\mathbf{w}\mathbf{x}}}$$

- This is equivalent to $\log \frac{p_1(\mathbf{x}; \mathbf{w})}{p_0(\mathbf{x}; \mathbf{w})} = \mathbf{w}\mathbf{x}$
- That is, the log odds of class 1 is a linear function of x
- Q: How to find **W**?

Logistic Regression

- Conditional data likelihood - Probability of observed Y values in the training data, conditioned on corresponding X values.
- We choose parameters w that satisfy

$$\mathbf{w} = \arg \max_{\mathbf{w}} \prod_l P(y^l | \mathbf{x}^l, \mathbf{w})$$

- where
 - $\mathbf{w} = \langle w_0, w_1, \dots, w_n \rangle$ is the vector of parameters to be estimated,
 - y^l denotes the observed value of Y in the l th training example, and
 - \mathbf{x}^l denotes the observed value of \mathbf{X} in the l th training example

Logistic Regression

- Equivalently, we can work with log of conditional likelihood:

$$\mathbf{w} = \arg \max_{\mathbf{w}} \sum_l \ln P(y^l | \mathbf{x}^l, \mathbf{w})$$

- Conditional data log likelihood, $l(\mathbf{W})$, can be written as

$$l(\mathbf{w}) = \sum_l y^l \ln P(y^l = 1 | \mathbf{x}^l, \mathbf{w}) + (1 - y^l) \ln P(y^l = 0 | \mathbf{x}^l, \mathbf{w})$$

- Note here that Y can take only values 0 or 1, so only one of the two terms in the expression will be non-zero for any given y^l

Logistic Regression: Training

- We need to estimate:

$$\mathbf{w} = \arg \max_{\mathbf{w}} \sum_l \ln P(y^l | \mathbf{x}^l, \mathbf{w})$$

- Equivalently, we can minimize negative log likelihood
- This is convex – so, unique global minimum
- No closed-form solution though. Iterative method required.

Logistic Regression: Training

- Use gradient ascent (descent) for the maximization (min) problem
- The i th component of the vector gradient has the form

$$\frac{\partial}{\partial w_i} l(\mathbf{w}) = \sum_l x_i^l (y^l - \hat{P}(y^l = 1 | \mathbf{x}^l, \mathbf{w}))$$

Logistic Regression prediction

- Beginning with initial weights, we repeatedly update the weights in the direction of the gradient, changing the i th weight according to

$$w_i \leftarrow w_i + \eta \sum_l x_i^l (y^l - \hat{P}(y^l = 1 | \mathbf{x}^l, \mathbf{w}))$$

Regularization in Logistic Regression

- Overfitting can arise especially when data has very high dimensions and is sparse.
- One approach -> modified “penalized log likelihood function,” which penalizes large values of \mathbf{w} , as before.

$$\mathbf{w} = \arg \max_{\mathbf{w}} \sum_l \ln P(y^l | \mathbf{x}^l, \mathbf{w}) - \frac{\lambda}{2} \|\mathbf{w}\|^2$$

- Derivative then becomes:

$$\frac{\partial}{\partial w_i} l(\mathbf{w}) = \sum_l x_i^l (y^l - \hat{P}(y^l = 1 | \mathbf{x}^l, \mathbf{w})) - \lambda w_i$$

Logistic Regression

- In general, NB and LR make different assumptions
 - NB: Features independent given class -> assumption on $P(X|Y)$
 - LR: Functional form of $P(Y|X)$, no assumption on $P(X|Y)$
- LR is a linear classifier
 - decision rule is a hyperplane
- LR optimized by conditional likelihood
 - no closed-form solution
 - Concave (convex) -> global optimum with gradient ascent (descent)
- Extending logistic regression to multiple classes
 - Use softmax for each class $k!$

$$p(y = k|x) = \frac{\exp(\theta_k^T x)}{\sum_{i=1}^K \exp(\theta_i^T x)}$$

Readings

- PRML Bishop, Chapter 4 (Sec 4.3)
- [“Introduction to Machine Learning” by Ethem Alpaydin](#), Chapter 10